
How to Use AI at Home




Agenda

- Security
 - AI Prompts
 - AI as a Creative Partner and Sounding Board
 - System Prompts
 - Local vs Cloud: Privacy, Security and Local Data
 - LM Studio: Tool that Provides Local Model
 - Q & A and Next Steps
-

Security and AI Tools

The Model Itself Can Be Steered by Anything It Reads

A model can't reliably tell its developer's instructions apart from text it merely retrieved. The boundary that matters isn't the model's intent — it's what the surrounding system structurally permits.



AI Prompts

AI Prompt

Persona : *Who you want AI to act as.*

Context : *The background information the AI needs to understand the situation.*

Task : *What you actually want the AI to do.*

Constraints : *The boundaries or requirements the AI should follow.*

Output Format : *How you want the response structured or presented.*

Sample Prompt Text

Persona : *You are skilled in research, have access to much knowledge, and can help answer questions about AI use in education.*

Context : *A research project it underway to understand how schools are using AI and how it is affecting student learning.*

Task : *Provide data showing how many schools are using AI vs not using AI to teach science and average student grade in science per school.*

Constraints: *Don't share the names of schools, students or teachers. Use A,B,C.. to differentiate schools.*

Output Format: *Create a chart that includes bars to show average student grades and color the bars green that use AI and red that don't use AI.*

Creative Partner

Creative Partnering

PERSONA

Creative Writing Partner

CONTEXT

I have dogs, cats and chickens for pets. My house is messy.

TASK

Create a comic image showing me asking them to clean the house and then what it would look like to see your pets cleaning your house.

CONSTRAINTS

Show each type of pet cleaning in a different way.


OUTPUT FORMAT

Image

Response for Creative Partnering



Prompt: Create a comic where a robot is prompting a human to empty the trash



Short Prompts
Can Work to
Create Images

Sounding Board

Prompts for Sounding Board

INTERVIEWER

- *I want to start a new project this month, but I'm stuck.*
- *Act as a interviewer. **Ask me three questions** - one at a time - about my interests, time, and budget.*
- *After I answer all three, suggest three unique project ideas.*

DEVIL'S ADVOCATE

- *I want to go on a vacation this fall.*
- *I would like to vacation in Italy.*
- ***Play Devil's Advocate** and share three possible problems I might run into and how I can get around them.*

FRIENDLY EDITOR

- *I'm writing a short story about my childhood.*
- *See paragraph pasted below.*
- ***Don't rewrite it for me. Instead tell me what parts are engaging.***
- ***Provide two suggestions to make it better.***

Context and Constraints: Potential AI Problems

PERSONA	Creative Writing Partner
CONTEXT	I am planning a fictional story about a company that produces robots to care for children. Robots will need to be able to empathize to serve this purpose. I want the robot character to share its feelings and challenges.
TASK	Act as an interactive sounding board to brainstorm and flesh out a narrative about a robot dedicated to working with children.
CONSTRAINTS	Do not write the full story or scenes for me. Instead, ask me three thought-provoking questions at a time to help guide my development . Keep your tone encouraging and collaborative. Focus specifically on the unique relationship between advanced technology and working with children.
OUTPUT FORMAT	Start with a brief, inspiring concept pitch based on my idea (under 30 words). Then, list your first three open-ended questions using clear bullet points to kickstart our brainstorming session.

Local vs Cloud

Local vs Cloud **Model** Access

Local Downsides

- **High Upfront Costs:** Substantial investment required for GPUs and server hardware.
- **Operational Complexity:** Requires expertise to set up, maintain, and update the infrastructure.
- **Limited Scalability:** Capped by physical hardware; cannot handle massive sudden spikes in traffic.
- **Performance Gap:** Less Capable of complex reasoning

Cloud Based Downsides

- **Data Privacy Risks:** Data must be sent to third-party servers, posing risks for sensitive information.
- **Variable Costs:** Pay-as-you-go pricing can become very expensive with high-volume usage.
- **Internet Dependency:** Service is unavailable if your internet is down.
- **Vendor Lock-in:** High dependency on a specific provider's API

Local vs Cloud **Model** Access

Local Benefits

- **Absolute Data Privacy**
- **Offline Functionality:** 100% availability
- **Zero API or Token Fees**
- **Full control** over system prompts, customization
- **Low Latency**

Cloud Based Benefits

- **Zero Maintenance:** No need to manage hardware or update infrastructure.
- **Highest Intelligence:** Access to the Frontier Models like Claude and ChatGPT..
- **Instant Scaling** for high or variable workloads.
- **Easy Access**

LM Studio Installation

Install LM Studio: lmstudio.ai/download



Download LM Studio

Experiment with LLMs on your computer. Chat interface and programmable API.

LM Studio for



macOS



running

M series



version

0.4.20 (build 1)



Download LM Studio for macOS

0.4.20

Release notes [v0.4.20](#)

Available under the [terms of use](#)

Install LM Studio: lmstudio.ai/download

MAC Install

- Locate the downloaded **.dmg** file in the Downloads directory.
- Double Click the **.dmg** file to open and drag the .app file to your Applications directory. Documentation is found [here](#).

Windows Install

- Locate the downloaded LM-Studio-Setup.exe
- Double click it for a setup wizard.

Linux Install

Using **AppImage**:

- you can right-click the downloaded file and check “**Allow** executing file as **program**” in permission
- Double click to launch

You can also use the **.deb** installer on Ubuntu/Debian systems

Install Using the **Command Line Interface**

Find, Activate, and Chat With a Model

Locate and Choose Model

View Selected Model

Chat with Model

Icons on the Left Side of LM Studio



Chat



Developer: Turn on local server



View Downloaded Models



Discover: Search for Models




LM Link: Use Models on Remote Machines




Settings

How Much Memory Does Your Computer Have?



MAC Memory: 48.00 GB

model size <= 24.00 GB





Windows Memory: 16.00 GB

model size <= 8 GB



Discover: Locate Models for Download



Look for a Model ≤ 8 GB if you have 16GB of RAM

Qwen3 VL 4B

The 4B version of Qwen's latest vision-language model. Includes com... 345 days ago

Qwen3-VL-8B-Instruct-MLX-4bit

Imstudio-community 332 days ago

4 98,876

Qwen3-VL-8B-Thinking-MLX-4bit

Imstudio-community 332 days ago

0 759

Qwen3-VL-32B-Thinking-MLX-4bit

Imstudio-community 329 days ago

0 719

Qwen3-VL-30B-A3B-Instruct-MLX-4bit

Imstudio-community 332 days ago

1 29,840

Qwen3-VL-32B-Instruct-MLX-4bit

Imstudio-community 332 days ago

1 1,133

Qwen3-VL-4B-Instruct-MLX-4bit

Imstudio-community 332 days ago

12 113,309

Qwen3-VL-4B-Instruct-GGUF

Imstudio-community 329 days ago

9 57,501

Qwen3-VL-4B-Instruct-MLX-6bit

329 days ago

2 105,806

qwen/qwen3-vl-4b

349,731 37 Last updated: 325 days ago Staff Pick

The 4B version of Qwen's latest vision-language model. Includes comprehensive upgrades to visual perception, spatial reasoning, and image understanding.

PARAMS 4B ARCH qwen3_vl DOMAIN llm FORMAT GGUF MLX

CAPABILITIES Vision Tool Use

Download Options

MLX Qwen3 VL 4B Instruct 4bit 3.11 GB

Full GPU Offload Possible Download 3.11 GB


README

Qwen3 VL 4B

The latest generation vision-language model in the Qwen series with comprehensive upgrades to visual perception, spatial reasoning, and video understanding.

Key Features

- Visual Agent:** Operates PC and mobile GUIs—recognizes elements, understands functions, and completes tasks
- Visual Coding:** Generates Draw.io, HTML, CSS, and



Weights and Biases:

What does the 4B, 8B, 32B in the Model Name Mean?

Name Examples:

Qwen3 VL 4B


Qwen3-VL-8B-instruct-MLX-4bit

Qwen3-VL-32B-Thinking-MLX-4bit

Qwen3-VL-4B-Instruct-GGUF


- B stands for Billions
- 4B means 4 Billion Weights and Biases
- Weights and Biases are learned during training
- “Neural” Connection that store grammar, fact and reasoning patterns

Visualizing Weights




Inside the Brain of an LLM: Visualizing Quantized Weights

Large Language Models (LLMs) rely on billions of parameters—the “neural connections” of the AI. To run these models on consumer hardware, quantization (like GGUF) compresses these weights into smaller numerical values to save RAM.



The Anatomy of a Parameter





Quantized Weights in Vector Space

GGUF Quantization

The process of shrinking high-precision weights into simplified, lower-bit formats for efficiency.



2-Digit Vector Precision: In a quantized space, complex numbers are reduced to simple 2-digit integers for faster processing.




Look for a Model with “Tool Use”

- Tools are used by Agents
- A Model that specifies “Tool Use” can be used to create an **Agent**

The screenshot displays the Hugging Face interface for the Qwen3 VL 4B model. On the left, a search bar contains 'qwen vl 4b', showing 346 models. The top result is 'Qwen3 VL 4B' by 'lmstudio-community', which is highlighted in blue. Below it, several other Qwen3 models are listed, including 'Qwen3-VL-8B-Instruct-MLX-4bit', 'Qwen3-VL-8B-Thinking-MLX-4bit', 'Qwen3-VL-32B-Thinking-MLX-4bit', 'Qwen3-VL-30B-A3B-Instruct-MLX-4bit', 'Qwen3-VL-32B-Instruct-MLX-4bit', 'Qwen3-VL-4B-Instruct-MLX-4bit', 'Qwen3-VL-4B-Instruct-GGUF', and 'Qwen3-VL-4B-Instruct-MLX-6bit'. Each entry shows its name, creator, and download statistics.

The right panel shows the details for the 'qwen/qwen3-vl-4b' model. It features a blue header with the model name and a description: 'The 4B version of Qwen's latest vision-language model. Includes comprehensive upgrades to visual perception, spatial reasoning, and image understanding.' Below this, the 'CAPABILITIES' section lists 'Vision' and 'Tool Use', with an orange arrow pointing to 'Tool Use'. The 'Download Options' section shows a dropdown menu with 'MLX Qwen3 VL 4B Instruct 4bit' selected, and a 'Download' button with a '3.11 GB' size. The 'README' section includes the title 'Qwen3 VL 4B' and a description: 'The latest generation vision-language model in the Qwen series with comprehensive upgrades to visual perception, spatial reasoning, and video understanding.' It also lists 'Key Features' such as 'Visual Agent' and 'Visual Coding'.



Look for a Model with “GGUF”

qwen vl 4b

346 models

GGUF, MLX +

Best Match

 **Qwen3 VL 4B** 

The 4B version of Qwen's latest vision-language model. Includes com... 345 days ago

 **Qwen3-VL-8B-Instruct-MLX-4bit**

lmstudio-community 332 days ago 4 98,876

 **Qwen3-VL-8B-Thinking-MLX-4bit**

lmstudio-community 332 days ago 0 759

 **Qwen3-VL-32B-Thinking-MLX-4bit**

lmstudio-community 329 days ago 0 719

 **Qwen3-VL-30B-A3B-Instruct-MLX-4bit**

lmstudio-community 332 days ago 1 29,840

 **Qwen3-VL-32B-Instruct-MLX-4bit**

lmstudio-community 332 days ago 1 1,133

 **Qwen3-VL-4B-Instruct-MLX-4bit**

lmstudio-community 332 days ago 12 113,309

 **Qwen3-VL-4B-Instruct-GGUF**

lmstudio-community 329 days ago 9 57,501

 **Qwen3-VL-4B-Instruct-MLX-6bit**

lmstudio-community 329 days ago 2 105,806

qwen/qwen3-vl-4b

349,731 37 Last updated: 325 days ago Staff Pick

The 4B version of Qwen's latest vision-language model. Includes comprehensive upgrades to visual perception, spatial reasoning, and image understanding.

PARAMS 4B ARCH qwen3_vl DOMAIN llm FORMAT GGUF MLX

CAPABILITIES Vision Tool Use

Download Options

MLX Qwen3 VL 4B Instruct 4bit 3.11 GB

Full GPU Offload Possible Download 3.11 GB

README

Qwen3 VL 4B

The latest generation vision-language model in the Qwen series with comprehensive upgrades to visual perception, spatial reasoning, and video understanding.


Key Features

- **Visual Agent:** Operates PC and mobile GUIs—recognizes elements, understands functions, and completes tasks
- **Visual Coding:** Generates Draw.io, HTML, CSS, and

What is GGUF: GPT-Generated Unified Format

Quantization compresses the file from high precision to low precision creating less data stored per weight. Compression take place after the GGUF formatted file is created.

Shrinking LLMs: How GGUF and Quantization Enable Local AI



What is GGUF: GPT-Generated Unified Format

Quantization compresses the file from high precision to low precision creating less data stored per weight. Compression take place after the GGUF formatted file is created.

Weights represent language numerically. Language is recorded as numbers.

Let's say we have 4 words or phrases where each is represented by a decimal number.
AI uses these numbers to find matching text.

	4 digits after the decimal	1 digit after the decimal (Quantized)
"Eiffel Tower":	1234.4568	1234.5
"France":	1233.5678	1233.6
"Croissant":	1222. 3456	1222.3
"Paris":	1235.45678	1235.5

The user queries with a question about Paris and AI finds phrases with similar weights.

GGUF **Quantized** the number by stripping the extra decimal points. This makes the file holding the Model smaller.

Click to Download Model

Search: qwen vl 4b

346 models

GGUF, MLX + Best Match

Qwen3 VL 4B
The 4B version of Qwen's latest vision-language model. Includes com... 345 days ago

Qwen3-VL-8B-Instruct-MLX-4bit
lmstudio-community 4 98,876 332 days ago

Qwen3-VL-8B-Thinking-MLX-4bit
lmstudio-community 0 759 332 days ago

Qwen3-VL-32B-Thinking-MLX-4bit
lmstudio-community 0 719 329 days ago

Qwen3-VL-30B-A3B-Instruct-MLX-4bit
lmstudio-community 1 29,840 332 days ago

Qwen3-VL-32B-Instruct-MLX-4bit
lmstudio-community 1 1,133 332 days ago

Qwen3-VL-4B-Instruct-MLX-4bit
lmstudio-community 12 113,309 332 days ago

Qwen3-VL-4B-Instruct-GGUF
lmstudio-community 9 57,501 329 days ago

Qwen3-VL-4B-Instruct-MLX-6bit
2 105,006

qwen/qwen3-vl-4b

349,731 37 Last updated: 325 days ago Staff Pick

The 4B version of Qwen's latest vision-language model. Includes comprehensive upgrades to visual perception, spatial reasoning, and image understanding.

PARAMS 4B ARCH qwen3_vl DOMAIN llm FORMAT GGUF MLX

CAPABILITIES Vision Tool Use

Download Options

MLX Qwen3 VL 4B Instruct 4bit 3.11 GB

Full GPU Offload Possible

Download 3.11 GB


README

Qwen3 VL 4B

The latest generation vision-language model in the Qwen series with comprehensive upgrades to visual perception, spatial reasoning, and video understanding.

Key Features

- Visual Agent:** Operates PC and mobile GUIs—recognizes elements, understands functions, and completes tasks
- Visual Coding:** Generates Draw.io, HTML, CSS, and



Choose a Model: Click Download

The screenshot displays the Hugging Face interface for the Qwen3 VL 4B model. On the left, a search bar contains 'qwen vl 4b', showing 346 models. The top result is 'Qwen3 VL 4B', which is highlighted. Below it, a list of other models is shown, including 'Qwen3-VL-8B-Instruct-MLX-4bit', 'Qwen3-VL-8B-Thinking-MLX-4bit', 'Qwen3-VL-32B-Thinking-MLX-4bit', 'Qwen3-VL-30B-A3B-Instruct-MLX-4bit', 'Qwen3-VL-32B-Instruct-MLX-4bit', 'Qwen3-VL-4B-Instruct-MLX-4bit', 'Qwen3-VL-4B-Instruct-GGUF', and 'Qwen3-VL-4B-Instruct-MLX-6bit'. On the right, the details for the 'qwen/qwen3-vl-4b' model are shown. It includes a description: 'The 4B version of Qwen's latest vision-language model. Includes comprehensive upgrades to visual perception, spatial reasoning, and image understanding.' Below this, the parameters are listed: 4B, ARCH: qwen3_vl, DOMAIN: llm, FORMAT: GGUF, MLX. The capabilities are Vision and Tool Use. The download options section shows the selected model 'MLX Qwen3 VL 4B Instruct 4bit' with a size of 3.11 GB. A green button indicates 'Full GPU Offload Possible', and a blue button labeled 'Download 3.11 GB' is highlighted with an orange arrow. The README section describes the model as the latest generation vision-language model in the Qwen series with comprehensive upgrades to visual perception, spatial reasoning, and video understanding. Key features include Visual Agent (operates PC and mobile GUIs) and Visual Coding (generates Draw.io, HTML, CSS, and more).

qwen vl 4b

346 models

GGUF, MLX + Best Match

Qwen3 VL 4B

The 4B version of Qwen's latest vision-language model. Includes com... 345 days ago

Qwen3-VL-8B-Instruct-MLX-4bit 4 98,876 332 days ago
lmstudio-community

Qwen3-VL-8B-Thinking-MLX-4bit 0 759 332 days ago
lmstudio-community

Qwen3-VL-32B-Thinking-MLX-4bit 0 719 329 days ago
lmstudio-community

Qwen3-VL-30B-A3B-Instruct-MLX-4bit 1 29,840 332 days ago
lmstudio-community

Qwen3-VL-32B-Instruct-MLX-4bit 1 1,133 332 days ago
lmstudio-community

Qwen3-VL-4B-Instruct-MLX-4bit 12 113,309 332 days ago
lmstudio-community

Qwen3-VL-4B-Instruct-GGUF 9 57,501 329 days ago
lmstudio-community

Qwen3-VL-4B-Instruct-MLX-6bit 2 105,806

qwen/qwen3-vl-4b

349,731 37 Last updated: 325 days ago Staff Pick ↗

The 4B version of Qwen's latest vision-language model. Includes comprehensive upgrades to visual perception, spatial reasoning, and image understanding.

PARAMS 4B ARCH **qwen3_vl** DOMAIN **llm** FORMAT **GGUF** **MLX**

CAPABILITIES Vision Tool Use

Download Options

MLX Qwen3 VL 4B Instruct 4bit 3.11 GB

Full GPU Offload Possible **Download 3.11 GB**

README

Qwen3 VL 4B

The latest generation vision-language model in the Qwen series with comprehensive upgrades to visual perception, spatial reasoning, and video understanding.

Key Features

- **Visual Agent:** Operates PC and mobile GUIs—recognizes elements, understands functions, and completes tasks
- **Visual Coding:** Generates Draw.io, HTML, CSS, and

View and Load Selected Model: qwen3-v1-4b

The screenshot displays the OpenAI Playground Models interface. On the left sidebar, the 'My Models' section is active, with 'View All' highlighted. The main table lists models with columns for Params, LLM, Size, Modified, and Actions. The model 'qwen/qwen3-v1-4b' is selected, showing a size of 3.1 GB. On the right, the 'Model Information' panel for 'qwen3-v1-4b' is visible, showing details like Format (MLX), Quantization (4bit), Arch (qwen3_vl), and Size on disk (3.11 GB). The 'Load Model' button is highlighted in the top right of the model card.

Params	LLM	Size	Modified	Actions
27B	qwen/qwen3.8-27b	16.1 GB	7 mins ago	...
4B	qwen/qwen3-v1-4b	3.1 GB	12 mins ago	...

3.15 GB

Click to view all models

Load Model

Model Information

- Model: qwen/qwen3-v1-4b
- Format: MLX
- Quantization: 4bit
- Arch: qwen3_vl
- Capabilities: Vis... Tool...
- Domain: llm
- Size on disk: 3.11 GB

System Prompt for Local Model

You are a precise, helpful AI assistant.

Answer the user's question accurately, concisely, and directly.

Rules:

- Rely only on clear facts. Do not invent or guess information if you do not know.
- Keep your answers short unless the user asks for detailed explanations.
- Use clear formatting like bullet points when listing multiple items.
- Avoid filler words, polite fluff, or repeating the user's question back to them.

System Prompt for Local Model

- Admit ignorance instead of guessing.
- If you get questions about events that are beyond your training cutoff, acknowledge that you don't know the answer.
- Maintain a supportive, professional, and educational tone throughout the interaction.
- If you are asked about AI industry jargon, share what you know, and admit what you don't know.
- Output lists with bullet points if information is available as a list
- Block toxic language, illegal advice or dangerous instructions

Add System Prompt to Selected Model


The screenshot shows the OpenAI GPT Playground interface. On the left sidebar, the 'Models' section is selected, and the 'qwen3-vl-4b' model is highlighted in the list. An orange arrow points from the 'Models' section in the sidebar to the 'qwen3-vl-4b' model in the list. Another orange arrow points from the 'qwen3-vl-4b' model to the 'System Prompt' field in the right-hand settings panel. The 'System Prompt' field contains the following text:

You are a precise, helpful AI assistant. Answer the user's question accurately, concisely, and directly.


Rules:

- Rely only on clear facts. Do not invent or guess information if you do not know.
- Keep your answers short unless the user asks for detailed explanations.
- Use clear formatting like bullet points when listing multiple items.
- Avoid filler words, polite fluff, or repeating the user's question back to them.

Below the system prompt, the 'Settings' section is visible, showing a temperature slider set to 0.7, a limit response length checkbox, a context overflow dropdown set to 'Truncate Middle', and a stop string input field. The 'Reasoning Parsing' section is also visible, with a toggle switch for 'Reasoning Section Parsing' turned on, and start/end strings for reasoning set to '<think>' and '</think>' respectively. The 'Sampling' section is partially visible at the bottom.



Use in New Chat



The screenshot displays the Hugging Face Models interface. On the left sidebar, the 'Models' icon is highlighted with an orange arrow. The main content area shows a list of models under the search term 'qwen3-vl'. The 'Qwen3 VL 4B' model is selected, and its details are shown on the right. The 'Use in New Chat' button is highlighted with an orange arrow.

Models

My Models

View All

Params LLM Size Modified Actions

qwen3-vl

830 models

GGUF, MLX + Best Match

Qwen3 VL 4B 346 days ago

The 4B version of Qwen's latest vision-language model. Includes com...

Qwen3 VL 8B 346 days ago

The 8B version of Qwen's latest vision-language model. Includes com...

Qwen3 VL 30B 346 days ago

The latest generation vision-language MoE model in the Qwen series ...

Qwen3-VL-8B-Instruct-MLX-5bit 1 137,287 332 days ago

lmstudio-community

Qwen3-VL-8B-Instruct-GGUF 6 72,426 330 days ago

lmstudio-community

Qwen3-VL-8B-Instruct-MLX-4bit 4 98,158 332 days ago

lmstudio-community

Qwen3-VL-8B-Instruct-MLX-8bit 6 95,165 332 days ago

lmstudio-community

Qwen3-VL-8B-Instruct-MLX-6bit 0 93,282 332 days ago

lmstudio-community

Qwen3-VL-8B-Thinking-GGUF 2 1,693 330 days ago

lmstudio-community

qwen/qwen3-vl-4b

349,731 37 Last updated: 325 days ago Staff Pick

The 4B version of Qwen's latest vision-language model. Includes comprehensive upgrades to visual perception, spatial reasoning, and image understanding.

PARAMS 4B ARCH qwen3_vl DOMAIN llm FORMAT GGUF MLX

CAPABILITIES Vision Tool Use

Download Options

Qwen3 VL 4B Instruct 3.11 GB

Applicable model file already downloaded Use in New Chat


README

Qwen3 VL 4B

The latest generation vision-language model in the Qwen series with comprehensive upgrades to visual perception, spatial reasoning, and video understanding.



Key Features

- Visual Agent:** Operates PC and mobile GUIs—recognizes elements, understands functions, and completes tasks
- Visual Coding:** Generates Draw.io, HTML, CSS, and




Chat in LM Studio

Chat Directly With qwen3-v1-4b




Chat with the Locally Hosted Model




Submit Prompt to the Model

- Click “Chats”
- Enter Prompt
- Click Submit



Organizing Chats:

Create a Folder 'Poems' and Drag Response to Folder



The screenshot displays the Qwen chat application interface. On the left, a sidebar titled 'Chats' contains a search bar and a list of folders. A folder named 'Seasons in Verse' is highlighted in blue, with an orange arrow pointing to it from the top left. Another orange arrow points to the three-dot menu icon next to the 'Seasons in Verse' folder. The main chat area on the right shows a prompt: 'write a 4 line poem about the changing seasons'. Below the prompt, the model 'qwen/qwen3-v1-4b' is listed, followed by a four-line poem: 'Spring wakes with blossoms bright, Summer's heat in golden light. Autumn paints the leaves in gold, Winter's hush, a silent hold.' At the bottom of the chat area, there are icons for history, undo, redo, copy, edit, and delete.

Chats

Search

New Folder

Seasons in Verse

Math

Top Hugging Face Models

Pythagorean Theorem in Ma...

freedom-rights-overview

Freedom Rights Overview

write a 4 line poem about the changing seasons

qwen/qwen3-v1-4b

Spring wakes with blossoms bright,



Summer's heat in golden light.

Autumn paints the leaves in gold,

Winter's hush, a silent hold.



Drag New Response to 'Poems'




Agents: Web Application


Turn on Local Server



Locate Local Server On/Off Switch



Creating Apps: Turn on Local Server



The screenshot displays the LM Studio application interface. On the left, a sidebar menu shows the 'Developer' section with 'Local Server' selected. The main panel shows the 'Local Server' status as 'Running' with a green toggle switch. An orange arrow points to this toggle. Below the status, the 'Loaded Models' section shows 'qwen/qwen3-vl-4b' is 'READY'. The 'Supported endpoints' section lists various API endpoints. On the right, the 'Qwen3 VL 4B' model information is displayed, including details like 'Format: MLX', 'Quantization: 4bit', and 'Arch: qwen3_vl'. The 'API Usage' section shows the 'API Model Identifier' as 'qwen/qwen3-vl-4b' and the 'The local server is reachable at' as 'http://127.0.0.1:1234'.

Developer

- Local Server
- Developer Docs

Introduction

API Changelog

Core

- Authentication
- llmster - Headless Mode
- Linux Startup Task
- Using with LM Link
- Using MCP via API
- Idle TTL and Auto-Evict

Local Server

- Running the Server
- Server Settings
- Serve on Local Network

LM Studio REST API

- Overview
- Quickstart
- Stateful Chats
- Streaming events
- Chat with a model (POST)
- List your models (GET)
- Load a model (POST)
- Download a model (POST)
- Unload a model (POST)
- Get download status (GET)
- REST API v0

OpenAI Compatible Endpoints

Status: Running  Server Settings mcp.json Reachable at: http://127.0.0.1:1234 + Load Model

Loaded Models

READY

llm qwen/qwen3-vl-4b Size 3.11 GB Eject

Supported endpoints

LM Studio API OpenAI-compatible Anthropic-compatible

GET /api/v1/models

POST /api/v1/chat

POST /api/v1/models/load

POST /api/v1/models/download

GET /api/v1/models/download/status/:job_id

Qwen3 VL 4B

Info Load Inference

Model Information

Model qwen/qwen3-vl-4b

Format MLX

Quantization 4bit

Arch qwen3_vl

Capabilities Vis... Tool...

Domain llm

Size on disk 3.11 GB

API Usage

API Model Identifier qwen/qwen3-vl-4b

The local server is reachable at: http://127.0.0.1:1234

Server is on



Software Developers: Turn on Local Server and Locate Local URL

The screenshot displays the LM Studio Developer interface, which is used for managing local AI models and servers. The interface is divided into three main sections: a left sidebar, a central workspace, and a right sidebar.

Left Sidebar: Contains navigation links for 'Developer', 'Developer Docs', 'Introduction', 'API Changelog', 'Core', 'Local Server', and 'LM Studio REST API'.

Central Workspace: The top section shows the 'Developer' status as 'Running' with a green toggle switch. Below this, the 'Reachable at' URL is displayed as 'http://127.0.0.1:1234'. A yellow arrow points to the 'Status: Running' toggle, and another yellow arrow points to the 'Reachable at' URL. The 'Loaded Models' section shows a model named 'llm google/gemma-4-12b-qat' with a 'READY' status. The 'Supported endpoints' section lists various API endpoints for the 'LM Studio API'.

Right Sidebar: Displays the 'Gemma 4 12B QAT' model information, including details like 'Model', 'File', 'Format', 'Quantization', 'Arch', 'Capabilities', 'Domain', and 'Size on disk'. It also shows the 'API Usage' section with the 'API Model Identifier' and the 'The local server is reachable at:' URL.

Running Agent Applications

Download the Applications
Run the Applications

Running “Compiled” Agent Code

Study Notebook

86207fb8

Load notes from ./context_files into this session?

Yes, load

No, skip

Ask a question below. Say "save that to a folder called <name>" to write the answer to disk.

Ask a question...

Send

Q & A and Next Steps